

[dotat.at](#) / [git](#) / [picoro.git](#) / blob[summary](#) | [shortlog](#) | [log](#) | [commit](#) | [commitdiff](#) | [tree](#)
[history](#) | [raw](#) | [HEAD](#)

commit ▼

? search:

☐ re

Remove C99 bool.

[\[picoro.git\]](#) / [picoro.c](#)

```
1 /*
2  * picoro - minimal coroutines for C.
3  * Written by Tony Finch <dot@dotat.at>
4  * http://creativecommons.org/publicdomain/zero/1.0/
5  */
6
7 #include <assert.h>
8 #include <setjmp.h>
9 #include <stdlib.h>
10
11 #include "picoro.h"
12
13 /*
14  * Each coroutine has a jmp_buf to hold its context when suspended.
15  *
16  * There are lists of running and idle coroutines.
17  *
18  * The coroutine at the head of the running list has the CPU, and all
19  * others are suspended inside resume(). The "first" coro object holds
20  * the context for the program's initial stack and also ensures that
21  * all externally-visible list elements have non-NULL next pointers.
22  * (The "first" coroutine isn't exposed to the caller.)
23  *
24  * The idle list contains coroutines that are suspended in
25  * coroutine_main(). After initialization it is never NULL except
26  * briefly while coroutine_main() forks a new idle coroutine.
27  */
28 static struct coro {
29     struct coro *next;
30     jmp_buf state;
31 } first, *running = &first, *idle;
32
33 /*
34  * A coroutine can be passed to resume() if
35  * it is not on the running or idle lists.
36  */
37 int resumable(coro c) {
38     return(c != NULL && c->next == NULL);
39 }
40
41 /*
42  * Add a coroutine to a list and return the previous head of the list.
43  */
44 static void push(coro *list, coro c) {
45     c->next = *list;
46     *list = c;
47 }
48
49 /*
50  * Remove a coroutine from a list and return it.
51  */
52 static coro pop(coro *list) {
53     coro c = *list;
54     *list = c->next;
55     c->next = NULL;
56     return(c);
57 }
58
59 /*
60  * Pass a value and control from one coroutine to another.
61  * The current coroutine's state is saved in "me" and the
62  * target coroutine is at the head of the "running" list.
```

```

63 */
64 static void *pass(coro me, void *arg) {
65     static void *saved;
66     saved = arg;
67     if(!setjmp(me->state))
68         longjmp(running->state, 1);
69     return(saved);
70 }
71
72 void *resume(coro c, void *arg) {
73     assert(resumable(c));
74     push(&running, c);
75     return(pass(c->next, arg));
76 }
77
78 void *yield(void *arg) {
79     return(pass(pop(&running), arg));
80 }
81
82 /* Declare for mutual recursion. */
83 void coroutine_start(void), coroutine_main(void*);
84
85 /*
86  * The coroutine constructor function.
87  *
88  * On the first invocation there are no idle coroutines, so fork the
89  * first one, which will immediately yield back to us after becoming
90  * idle. When there are idle coroutines, we pass one the function
91  * pointer and return the activated coroutine's address.
92  */
93 coro coroutine(void *fun(void *arg)) {
94     if(idle == NULL && !setjmp(running->state))
95         coroutine_start();
96     return(resume(pop(&idle), fun));
97 }
98
99 /*
100  * The main loop for a coroutine is responsible for managing the "idle" list.
101  *
102  * When we start the idle list is empty, so we put ourselves on it to
103  * ensure it remains non-NULL. Then we immediately suspend ourselves
104  * waiting for the first function we are to run. (The head of the
105  * running list is the coroutine that forked us.) We pass the stack
106  * pointer to prevent it from being optimised away. The first time we
107  * are called we will return to the fork in the coroutine()
108  * constructor function (above); on subsequent calls we will resume
109  * the parent coroutine_main(). In both cases the passed value is
110  * lost when pass() longjmp()s to the forking setjmp().
111  *
112  * When we are resumed, the idle list is empty again, so we fork
113  * another coroutine. When the child coroutine_main() passes control
114  * back to us, we drop into our main loop.
115  *
116  * We are now head of the running list with a function to call. We
117  * immediately yield a pointer to our context object so our creator
118  * can identify us. The creator can then resume us at which point we
119  * pass the argument to the function to start executing.
120  *
121  * When the function returns, we move ourselves from the running list to
122  * the idle list, before passing the result back to the resumer. (This
123  * is just like yield() except for adding the coroutine to the idle
124  * list.) We can then only be resumed by the coroutine() constructor
125  * function which will put us back on the running list and pass us a
126  * new function to call.
127  *
128  * We do not declare coroutine_main() static to try to stop it being inlined.
129  *
130  * The conversion between the function pointer and a void pointer is not
131  * allowed by ANSI C but we do it anyway.
132  */
133 void coroutine_main(void *ret) {
134     void *(*fun)(void *arg);
135     struct coro me;
136     push(&idle, &me);

```

```
137     fun = pass(&me, ret);
138     if(!setjmp(running->state))
139         coroutine_start();
140     for(;;) {
141         ret = fun(yield(&me));
142         push(&idle, pop(&running));
143         fun = pass(&me, ret);
144     }
145 }
146
147 /*
148  * Allocate space for the current stack to grow before creating the
149  * initial stack frame for the next coroutine.
150  */
151 void coroutine_start(void) {
152     char stack[16 * 1024];
153     coroutine_main(stack);
154 }
155
156 /* eof */
```

Various little pure C coroutine implementations.

[Atom](#) [RSS](#)